


PRIOR — SHARED, UNCHANGED ACROSS ALL PHASES



Phase 2 (not yet implemented; dashed). Audio enters only to derive per-note targets. The observation model, the prior and the conjugate posterior are unchanged — what changes is the channel set and that the observation noise Σ_e is no longer near-zero.